

Random Forest Classifiers

Project Guide

- [Project Overview](#)
- [Part 1: Acquire, Explore, and Preprocess Data](#)
- [Part 2: Trees and Forests](#)
- [Part 3: Parameter Tuning](#)

Project Overview

EXPECTED TIME: 2 HRS

This assignment uses a variety of decision tree based classifiers to attempt prediction of whether or not a customer will default on their loans. Our data comes from [the UCI machine learning Repository on default of credit card clients](#)

Activities will include:

- Manipulating DataFrames
- Visualizing data
- Calculation of impurity measures
- Using `sklearn`'s tree and forest models
- Evaluate the effects of hyperparameter tuning

Motivation: Decision Trees and Forests offer easy to understand yet fairly advanced models with a variety of hyper-parameters to increase or decrease complexity to tune between bias and variance

Problem:

Given a number of personal variables, (sex, education, marriage status, age); and recent payment history, attempt to predict whether or not a customer will default in the next month.

Data:[Defaults of credit card clients](#)

Please see above link for a complete description of the data.

Data Set Information:

This research aimed at the case of customers' default payments in Taiwan and compares the predictive accuracy of probability of default among six data mining methods. From the perspective of data mining, this study reviewed the literature and used the following 23 variables as explanatory variables.

Attribute Information:

This research employed a binary variable, default payment (Yes = 1, No = 0), as the response variable. This study reviewed the literature and used the following 23 variables as explanatory variables:
X1: Amount of the given credit (NT dollar): it includes both the individual consumer credit and his/her family (supplementary) credit.
X2: Gender (1 = male; 2 = female).
X3: Education (1 = graduate school; 2 = university; 3 = high school; 4 = others).
X4: Marital status (1 = married; 2 = single; 3 = others).
X5: Age (year).
X6 - X11: History of past payment. We tracked the past monthly payment records (from April to September, 2005) as follows: X6 = the repayment status in September, 2005; X7 = the repayment status in August, 2005; X8 = the repayment status in July, 2005; X9 = the repayment status in June, 2005; X10 = the repayment status in May, 2005; X11 = the repayment status in April, 2005.
X12-X17: Amount of bill statement (NT dollar). X12 = amount of bill statement in September, 2005; X13 = amount of bill statement in August, 2005; X14 = amount of bill statement in July, 2005; X15 = amount of bill statement in June, 2005; X16 = amount of bill statement in May, 2005; X17 = amount of bill statement in April, 2005.
X18-X23: Amount of previous payment (NT dollar). X18 = amount paid in September, 2005; X19 = amount paid in August, 2005; X20 = amount paid in July, 2005; X21 = amount paid in June, 2005; X22 = amount paid in May, 2005; X23 = amount paid in April, 2005.

Part 1: Acquire, Explore, and Preprocess Data

Import / Read in Data

```
# Import necessary packages
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np

#
%matplotlib inline

# Read in Data
df = pd.read_excel("../resource/asnlib/publicdata/default of credit card clients.xls", header = 1)

df.rename(columns = {"PAY_0": "PAY_1"}, inplace = True) #renaming mis-named column
df.head()
```

Explore the Data

Check for nulls

```
print("Data Shape: ", df.shape, "\n")
df.info()
```

Appears no null data.

Investigate Categorical Variables

Question 1:

The "default payment next month" will be the target used for classification. Investigate its distribution for the question below.

```
### GRADED

### Assign an int to `ans0` corresponding to the total number of non-default records
### Assign an int to `ans1` corresponding to the total number of default records

### Note the a "0" in the 'default payments next month' column means "non-default" and a 1 means "default"

### YOUR ANSWERS BELOW

ans0 = 23364
ans1 = 6636
```

Question 2:

```
### GRADED

### This proportion of default and non-default records means we are dealing with

### 'a') Balanced classes
### 'b') Unbalanced classes
### Assign the character associated with your choice as a string to ans1

### YOUR ANSWER BELOW

ans1 = 'b'
```

Investigating the Features:

```
df['EDUCATION'].value_counts()
```

Note: The code-book only describes the education variable as having four values (1-4), yet, here, there are seven values (0-6). In some cases this might be grounds to throw out these unknown values (0,5,6). For now, we will leave them in, assuming that they have some (unknown to us) meaning.

```
df['MARRIAGE'].value_counts()
```

Note: Again, the code book only describes three values for marriage (1-3), yet here, "0" also appears. Given what we saw above, we might assume the "0" in these categorical variables is functionally used for the "null" value

```
df['SEX'].value_counts()
```

Note: A slight imbalance exists representation of men and women, with women making up a little over 60% of our observations. Thankfully this column contains no "0"s.

A Closer look at the "PAY" variables

```
df['PAY_1'].value_counts()
```

Question 3:

```
### GRADED

### True or False:
### The code-book (description of variables that came with the data set; copied above and in link)
### described what a "0" in the 'Pay_#' columns means?

### Assign boolean answer to ans1
### YOUR ANSWER BELOW

ans1 = False
```

Investigate relationship between "Pay", "Bill_amt" and "Pay_amt" variables:

```
for i in [-2,-1,0,1,2,8]:
    print(df[df['PAY_1']==i][['PAY_1','BILL_AMT1','PAY_AMT1']].head(8), "\n")
```

Unclear exactly how these "PAY" variables work. Possibly they should be treated as categorical data instead of discrete and interval data, but, for this modeling task, keep them as interval data.

```
### Define function for creating histograms
def pay_hist(df, cols, ymax):
    plt.figure(figsize=(10,7)) # define fig size

    for index, col in enumerate(cols): # For each column passed to function
        plt.subplot(2,3, index + 1) # plot on new subplot
        plt.ylim(ymax = ymax) # standardize ymax
        plt.hist(df[col]) # create hist
        plt.title(col) # title with column names
    plt.tight_layout(); # make sure titles don't overlap
```

```
pay_cols = ["PAY_"+str(n) for n in range(1,7)]
pay_amt_cols = ["PAY_AMT" + str(n) for n in range(1,7)]
bill_amt_cols = ["BILL_AMT" + str(n) for n in range(1,7)]
```

```
pay_hist(df, pay_cols, 20000)
```

Note: Clearly the "0" is the majority class for all of the "PAY" variables. But, unclear what a "0" means as it is not in the code book.

```
pay_hist(df, pay_amt_cols, 20000)
```

```
df[pay_amt_cols].boxplot();
```

```
df_no_0_pay_amt_1 = df[df["PAY_AMT1"]!=0]
df_no_0_pay_amt_1["PAY_AMT1"].hist()
```

Even taking out all the PAY_AMT of 0, most payments stay close to 0 with a long tail.

```
log_pay_amt1 = np.log10(df_no_0_pay_amt_1["PAY_AMT1"])
plt.hist(log_pay_amt1)
plt.title("Log10-Transformed values for 'PAY_AMT1' (Excluding 0s)");
```

Log Transformation (used above) or $\sqrt[4]{x}$ transformations (which automatically deals with 0s) can be a good way of looking more closely at skewed data. Above, we see that most of the repayment amounts are in the 1000's of dollars.

```
pay_hist(df, bill_amt_cols, 23000)
```

```
df[bill_amt_cols].boxplot();
```

Preprocessing

Currently "Sex" is coded as 2 for "female" and 1 for "male". As this is encoded as binary we will change the column name to "FEMALE" and subtract 1 from each value - thus 1 will be "Female" and 0 "Male".

Both Education and Marriage are categorical with multiple options. `pd.get_dummies()` will allow us to create n-1 binary features to encode the n categories.

Note regarding `get_dummies()`

`pd.get_dummies()` is **NOT** appropriate in many ML applications.

For Example: Currently we have the values 0-6 in the "EDUCATION" feature. Suppose that when we split the training and test data, all 14 of the "0" values ended up in the test set. Running `pd.get_dummies()` on the test set would add 6 columns where the training set would only have 5 columns added from that action. then, whatever model had been fit would not know how to deal that new feature / category.

The function is used here for simplicities' sake

```
df['SEX'] = df['SEX']-1 # change vals of 'sex' to 0,1

df.rename(columns = {'SEX':'FEMALE', "default payment next month":"default"}, inplace = True) # rename col names

for col, pre in zip(["EDUCATION", "MARRIAGE"],["EDU","MAR"]): # get dummies and rename cols for ed and marriage
    df = pd.concat([
        df.drop(col, axis = "columns"), pd.get_dummies(df[col], prefix = pre, drop_first = True)],
        axis = 'columns')

df.head()
```

Question 4:

```
### GRADED
### Why is it inappropriate to use `pd.get_dummies()` when preprocessing data?

### 'a') it destroys data by removing categorical data
### 'b') it will skew results of models
### 'c') it cannot deal with "new" categories possibly found in a test set
### 'd') it is not a real method
### Assign character corresponding to your choice as string to ans1.

### YOUR ANSWER BELOW

ans1 = 'c'
```

Question 5

```
### GRADED
### Why are 'dummy' variables necessary?
```

```

### 'a') It helps to expand the number of features
### 'b') Models cannot deal multi-class categorical data
### 'c') It translates categorical data into quantitative data
### Assign character corresponding to your choice as string to ans1.

### YOUR ANSWER BELOW

ans1 = 'b'

```

Trees and Forests: Intro

Given a target (categorical or continuous), a decision tree iteratively splits data. It splits data at the value in whichever feature that creates the greatest separation among the target variable.


For example, say the target variable on a dataset of professional basketball players is league: NBA and WNBA. The available features are height and weight.

Split 1: Although there are WNBA players over 6'6" tall, that is very rare. Thus, a decision tree might split the data at the height of 6'8", 6'7" or 6'6". At that split the tree would predict that any basketball player over 6'6" plays in the NBA.

Split 2: For players shorter than 6'6", the tree might decide to split that data again at 5'10" and predict that basketball players under 5'10" play in the WNBA.

Split 3: For those players between 5'10" and 6'6", maybe the tree might discriminate on weight, predicting that all those players who weigh more than 180 Lbs. play in the NBA.

The resulting tree could be visualized as below



Depth

The above tree has a depth of 2.

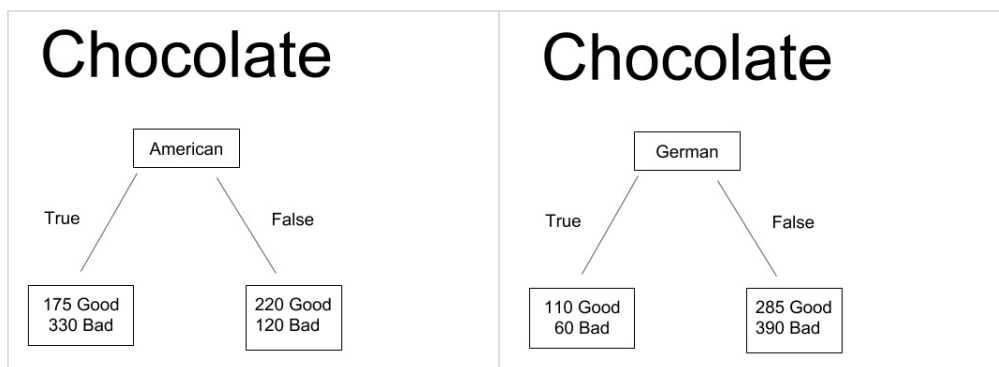
The maximum depth to which a tree is allowed to grow can be specified with `max_depth` in `sklearn`. By default, `max_depth` is set to `None` which means the tree will grow until all leaves (terminal nodes) are pure, or until other user-specified criteria are met. Importantly, `max_depth` can impact the amount of time it takes to build a tree (this becomes especially important when starting to work with forests.)

Splitting

The splits in the above trees were determined intuitively. Thankfully decision Trees do not make their decisions using intuition. In the `sklearn` package two splitting criteria are available for classifiers; "gini" and "entropy". In general, "gini" splitting favors larger partitions, where "entropy" favors splitting of smaller groups that are of a single class.

[More on Gini/Entropy](#)

Gini and Entropy Calculations



Above we have a split in a data set regarding good and bad chocolate.

You will be asked to calculate the gini impurity in the next questions.

The Split on "American" shows 175 good chocolates and 330 bad chocolates where `American == True`. There are 200 good and 120 bad chocolates where `American == False`.

The Split on "German" shows 110 good chocolates and 60 bad chocolates where `German == True`. There are 285 good and 390 bad chocolates where `German == False`.

As a reminder; the gini-index for a node is: $\sum_{j=1}^n p_j^2$

Where there are n classes and p_j is the frequency of class j in that node.

Finally, the indexes for each of these nodes is weighted by the proportion of data at each node, then summed.

Remember gini-indices closer to 0 are more "pure"

Question 6:

```
### GRADED
### Calculate the gini impurity for the split on 'American'.

### Answer will be tested to 4 decimal places
### Assign float answer to ans1

### YOUR ANSWER BELOW
def calcGI(g1, b1, g2, b2):
    t = sum([g1,b1,g2,b2])

    def calcNode(g,b, t):
        tn = g+b
        prop = ((g+b)/t)
        node = (1- ((g/tn)**2)- ((b/tn)**2))
        # print(prop, node)
        return prop*node

    return calcNode(g1,b1, t) + calcNode(g2,b2,t)

ans1 = calcGI(175, 330, 220, 120) #-> ~ .4544
```

Question 7:

```
### GRADED
### Same instructions as above, but this time, calculate gini impurity for split on 'German'
### Assign float answer to ans1

### YOUR ANSWER BELOW

def calcGI(g1, b1, g2, b2):
    t = sum([g1,b1,g2,b2])
    def calcNode(g,b, t):
        tn = g+b
        prop = ((g+b)/t)
        node = (1- ((g/tn)**2)- ((b/tn)**2))
        print(prop, node)
        return prop*node

    return calcNode(g1,b1, t) + calcNode(g2,b2,t)

ans1 = calcGI(110, 60, 285, 390) #-> 0.4816336
```

Question 8:

```
### GRADED
### According to the gini values calculated above, which split is better?

### 'a') American
### 'b') German
### Assign character associated with your choice as string to ans1.

### YOUR ANSWER BELOW

ans1 = 'a'
```

Forests

"Forests" are collections of decision trees designed to protect against over-fitting.

A single decision tree (particularly one that is allowed to grow to any depth), may be prone to overfitting. Algorithmically, the tree is designed to continue to make splits until it has completely classified all of the available data, and/or exhausted every possible split. Thus, no complexity and no particular is too fine for the tree to split upon. A tree might be "pruned" (by setting `max_depth`) to protect against over fitting, but, a "forest" of trees might also be used instead.

Building Tree / Forest Models

Question 9:

Baseline Accuracy

```
### GRADED
### Assume we are predicting default/not-default.
### What percent (int between 0 and 100) of observations would be correctly predicted if the majority class
### were predicted every time.
### E.g., What is the baseline accuracy?
### Assign int to ans1

### YOUR ANSWER BELOW

ans1=max(df['default'].value_counts()/df.shape[0])*100 # --> ~78
```

Import models and metrics; create train-test split

```
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, ExtraTreesClassifier, BaggingClassifier
```

```

from sklearn.metrics import classification_report
from sklearn.model_selection import train_test_split

# Create tts
X_train, X_test, y_train, y_test = train_test_split(
    df.drop("default", axis = 'columns'), df['default'],
    test_size = .3, random_state = 1738)

# Instantiate tree and forest models
dt = DecisionTreeClassifier()
bag = BaggingClassifier()
rf = RandomForestClassifier()
et = ExtraTreesClassifier()

```

Look at performance of classifiers using default parameters

```

dt.fit(X_train, y_train)
print("Decision Tree: \n", classification_report(y_test, dt.predict(X_test)), "\n")

print("-----")

bag.fit(X_train, y_train)
print("Bagging: \n", classification_report(y_test, bag.predict(X_test)), "\n")

print("-----")

rf.fit(X_train, y_train)
print("Random Forest: \n", classification_report(y_test, rf.predict(X_test)), "\n")

print("-----")

et.fit(X_train, y_train)
print("Extra Trees: \n", classification_report(y_test, et.predict(X_test)), "\n")

```

Question 10

Values of different predictions

```

### GRADED
### Are all predictions in this data set equal in business value?
### e.g. is correctly predicting a "default" as valuable as predicting a "non-default"
### 'a') Equal values
### 'b') Unequal values
### Assign string associated with your choice to ans1

### YOUR ANSWER BELOW

ans1 = 'b'

```

Part 3: Hyper Parameter Tuning

It appears both the Random Forest and the Extra Trees perform in a similarly effective manner. For this section, we will try to increase the performance of the Random Forest by tuning a couple hyper parameters, namely `criterion` and `estimators`

Because we are more interested in forecasting defaults than non-defaults, we will optimize on the recall of defaults -- recall is the proportion of defaults predicted over total defaults.

The below will not run on Vocareum due to processing constraints, thus the output is copied below :

```

# %%time
# from sklearn.metrics import recall_score
# criterion = ['gini', 'entropy']
# n_estimators = [5,10,20, 50, 100]
# scores = dict()
# i = 0
# for c in criterion:
#     for e in n_estimators:
#         rf = RandomForestClassifier(n_estimators = e, criterion = c, random_state = 1738)
#         rf.fit(X_train, y_train)
#         scores[i] = {'recall':recall_score(y_test, rf.predict(X_test)), 'trees' :e, "crit":c}
#         i+=1

```

```

# pd.DataFrame(scores).T

```

```
pd.DataFrame(scores).T
```

Out[12]:

		crit	recall	trees
0		gini	0.365679	5
1		gini	0.318693	10
2		gini	0.333504	20
3		gini	0.350868	50
4		gini	0.350868	100
5		entropy	0.364147	5
6		entropy	0.321757	10
7		entropy	0.331971	20
8		entropy	0.341164	50
9		entropy	0.342697	100

Question 11:

```
### GRADED
### According to the output shown in the picture above, the model with the best recall score featured:
### How many trees? Answer with int assigned to ans1.

### Which splitting method?
### 'a') entropy or 'b')gini.
### Put the character coresponding to your selection as a string in ans2

### YOUR ANSWER BELOW

ans1 = 5
ans2 = 'b'
```

With a general sense that maybe a `RandomForestClassifier` will perform best using `gini` splitting with somew here around 5 trees, further hyper-parameter tuning below.

```
# n_estimators = [1,2,3,4,5,6,7,8]
# scores2 = dict()
# i = 0
# for e in n_estimators:
#     rf = RandomForestClassifier(n_estimators = e, criterion = 'gini', random_state = 1738)
#     rf.fit(X_train, y_train)
#     scores2[i] = {'recall':recall_score(y_test, rf.predict(X_test)), 'trees' :e}
#     i+=1
```

```
# pd.DataFrame(scores2).T
```

```
pd.DataFrame(scores2).T
```

Out[14]:

	recall	trees
0	0.416241	1.0
1	0.223698	2.0
2	0.375894	3.0
3	0.267109	4.0
4	0.365679	5.0
5	0.281920	6.0
6	0.367722	7.0
7	0.305414	8.0

Question 12:

```
### GRADED
### On drilling down and focusing on 'gini' random forests with between 1 and 8 trees it becomes clear that
### hyper-parameter tuning here is creating:

### 'a') more consistency and better result
### 'b') more consistency and worse results
### 'c') no identifiable increase in consistency
### Assign the character corresponding to your choice as a string to ans1

### YOUR ANSWER BELOW

ans1 = 'c'
```

While it might be tempting to try to tune-and-tune-and-tune hyper-parameters to increase scores, in many cases (as in the example above) hyper-parameter tuning mostly results in *not* creating a better model, but a model that does a better job of predicting the test set.